

# Mitigating Backdoors/Trojans in Deep Neural Networks

CSIT375/975 AI and Cybersecurity

Dr Wei Zong

SCIT University of Wollongong

Disclaimer: The presentation materials come from various sources. For further information, check the references section

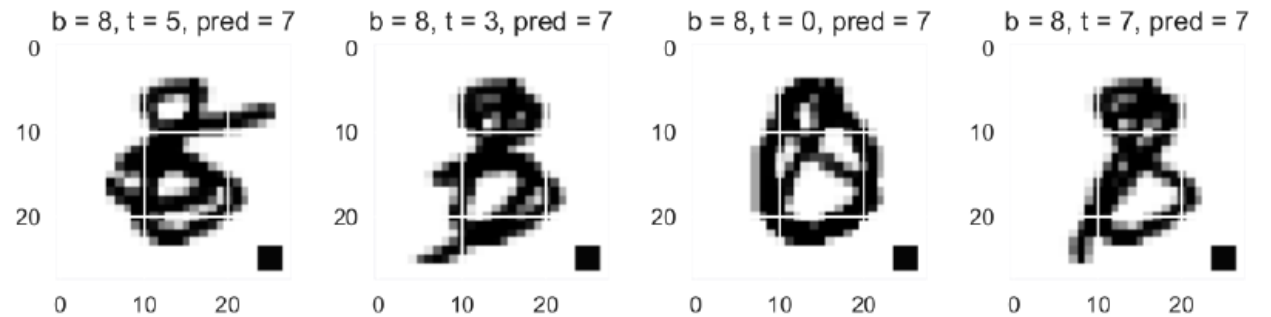
# Outline

- Detect triggers in input
  - STRIP
- Remove backdoors/Trojans in DNN
  - Fine-pruning
  - Neural cleanse
- Robust learning against backdoors/Trojans
  - Anti-backdoor learning

# Detect Triggers in Input

- Problem: can we detect whether input contains a trigger?
  - Backdoor may exist.
    - But not necessarily.
  - Do not have any information about triggers and target labels.
    - Adversaries will not share such information.
  - If a trigger is detected
    - Reject the input.

# Defense: STRIP



- Observation
  - Empirically, triggers are input-agnostic, e.g., BadNets.
    - Examples are shown in the figure on the top.
    - If a trigger exists, the output will be the same regardless the input content.
  - This inspires the strategy to detect Trojan attacks via repeatedly mixing input with another clean input which has a different label.
    - The intuition is that predictions for clean input will be altered randomly.
      - Input is ambiguous.
    - Predictions for input with triggers will stay stable.
  - Block input if triggers are identified in input.
    - No need to patch the model as malicious input are rejected.

# Defense: STRIP

- STRIP algorithm
  - An input is mixed with multiple other clean input to form a perturbed set.
    - Clean input are randomly drawn from the dataset.
  - Entropy of each perturbed input is then calculated.

$$\mathbb{H}_n = - \sum_{i=1}^M y_i \times \log_2 y_i$$

- where  $\mathbb{H}_n$  is entropy for the  $n^{th}$  perturbed input.
- $y_i$  indicates the probability of being classified as class  $i$ .
- $M$  is the total number of classes.
- Entropy value ranges  $[0, 1]$ 
  - A larger entropy means more randomness.
  - A smaller entropy means less randomness.

---

**Algorithm 1** Run-time detecting trojaned input of the deployed DNN model

---

```
1: procedure detection ( $x, \mathcal{D}_{test}, F_{\Theta}()$ , detection boundary )
2:    $trojanedFlag \leftarrow \text{No}$ 
3:   for  $n = 1 : N$  do
4:     randomly drawing the  $n_{th}$  image,  $x_n^t$ , from  $\mathcal{D}_{test}$ 
5:     produce the  $n_{th}$  perturbed images  $x^{p_n}$  by superimposing in-
       coming image  $x$  with  $x_n^t$ .
6:   end for
7:    $\mathbb{H} \leftarrow F_{\Theta}(\mathcal{D}_p)$   ▷  $\mathcal{D}_p$  is the set of perturbed images consisting of
        $\{x^{p_1}, \dots, x^{p_N}\}$ ,  $\mathbb{H}$  is the entropy of incoming input  $x$  assessed by
       averaging all the calculated entropy.
8:   if  $\mathbb{H} \leq \text{detection boundary}$  then
9:      $trojanedFlag \leftarrow \text{Yes}$ 
10:  end if
11:  return  $trojanedFlag$ 
12: end procedure
```

---

# Defense: STRIP

- STRIP algorithm (continued)
  - The entropy values are averaged
    - A **larger** entropy means higher possibility for the input being clean.
      - Perturbed input are ambiguous.
    - A **smaller** entropy means higher possibility for the existence of a trigger.
      - The trigger is detected.
  - Anomaly detection is employed to detect the existence of a trigger in new input.
    - Assume the entropy for clean input follows a **Normal (Gaussian) Distribution**,.
    - In practice, the entropy distribution for clean input can be calculated in advance to determine the detection threshold.

---

**Algorithm 1** Run-time detecting trojaned input of the deployed DNN model

---


```
1: procedure detection ( $x, \mathcal{D}_{test}, F_{\Theta}()$ , detection boundary )
2:    $trojanedFlag \leftarrow \text{No}$ 
3:   for  $n = 1 : N$  do
4:     randomly drawing the  $n_{th}$  image,  $x_n^t$ , from  $\mathcal{D}_{test}$ 
5:     produce the  $n_{th}$  perturbed images  $x^{p_n}$  by superimposing in-
       coming image  $x$  with  $x_n^t$ .
6:   end for
7:    $\mathbb{H} \leftarrow F_{\Theta}(\mathcal{D}_p)$   $\triangleright \mathcal{D}_p$  is the set of perturbed images consisting of
        $\{x^{p_1}, \dots, x^{p_N}\}$ ,  $\mathbb{H}$  is the entropy of incoming input  $x$  assessed by
       averaging all the calculated entropy.
8:   if  $\mathbb{H} \leq \text{detection boundary}$  then
9:      $trojanedFlag \leftarrow \text{Yes}$ 
10:  end if
11:  return  $trojanedFlag$ 
12: end procedure
```

---

# Defense: STRIP

- Results

- An example of entropy distribution for clean input and input with triggers is shown in the figure.
  - 2000 benign and 2000 Trojaned input images of GTSRB.
  - The entropy of input containing a trigger is concentrated at low values.
  - The entropy distribution for clean input spreads across a large range.
    - Consistent with the intuition that there is more randomness in predictions for clean input when mixed with other clean input.
  - The entropy distribution for clean input visually follows a normal distribution.
- Choosing a 1% false rejection rate (FRR) suppresses false acceptance rate (FAR) to be less than 1%.
  - Based on case studies on MNIST, CIFAR10, and GTSRB.
  - The FRR is the probability when the benign input is regarded as a trojaned input.
  - The FAR is the probability when the trojaned input is recognized as the benign input.




# Remove backdoors/Trojans in DNN

- Merely detecting trigger in input is not enough.
  - The risk does not disappear.
  - Backdoored models need to be purified.
- Problem: given a well trained DNN, can we remove **potential** backdoors?
  - Backdoor may exist.
    - But not necessarily.
  - Negligibly affect model performance.
    - Otherwise, decreasing its value.
  - A user may not have access to the original training set.
    - Download a pretrained model.
  - Do not have any information about triggers and target labels.
    - Adversaries will not share such information.
  - Computational costs need to be considered
    - Significantly less than training a clean model from scratch.



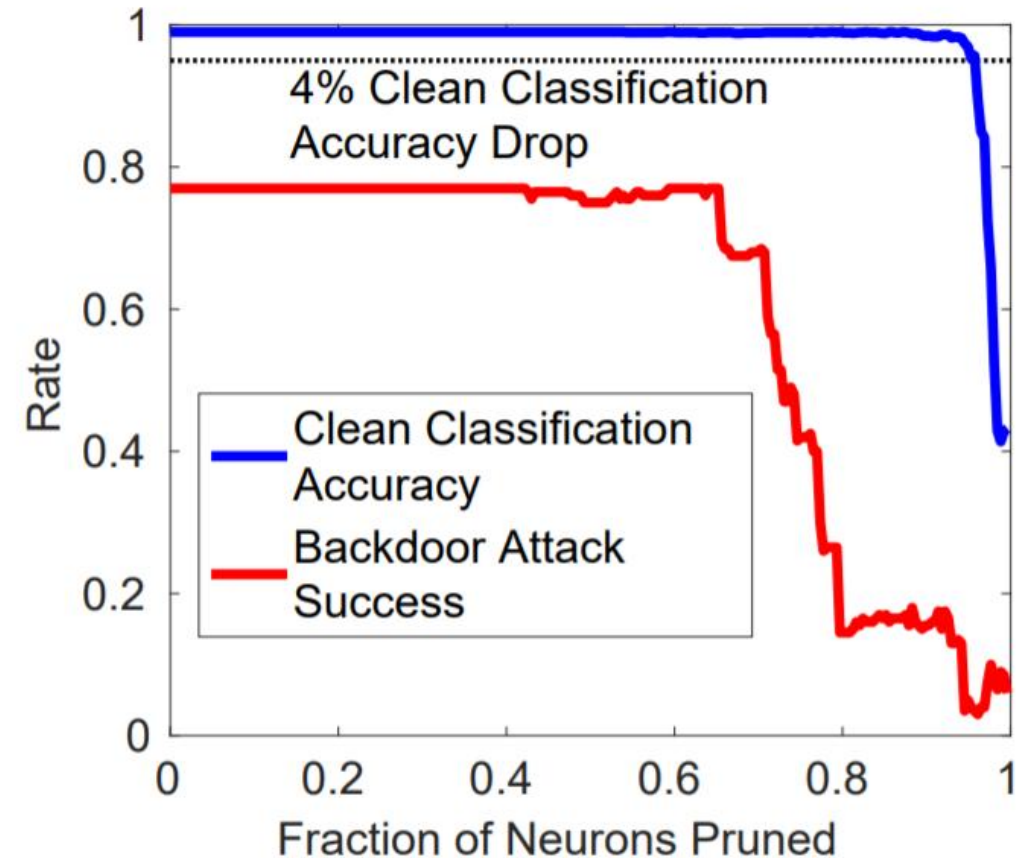
# Defense: Fine-Pruning

- Observation
  - Different neurons are activated for clean input and input with triggers.
    - A potential explanation is that each neuron aims to detect a specific feature from the input.
  - Hence, neurons that detect the existence of triggers are not activated when input is clean and vice versa.
- Key idea
  - Prune neurons that do not activate for clean input.
    - The purpose is to remove potential Trojan from a target model.



# Defense: Fine-Pruning

- A naïve approach fails
  - Initial results show that removing neurons that are not activated for normal input can remove potential Trojan.
    - This will also degrade the performance **significantly**.
    - This is because the original architecture is changed.
  - Trojan is successfully removed at the cost of 4% decrease in accuracy for clean input.
    - 4% decrease in accuracy is not negligible.
    - Research community improved the state-of-the-art top-1 accuracy on ImageNet by about only one percent point per year.

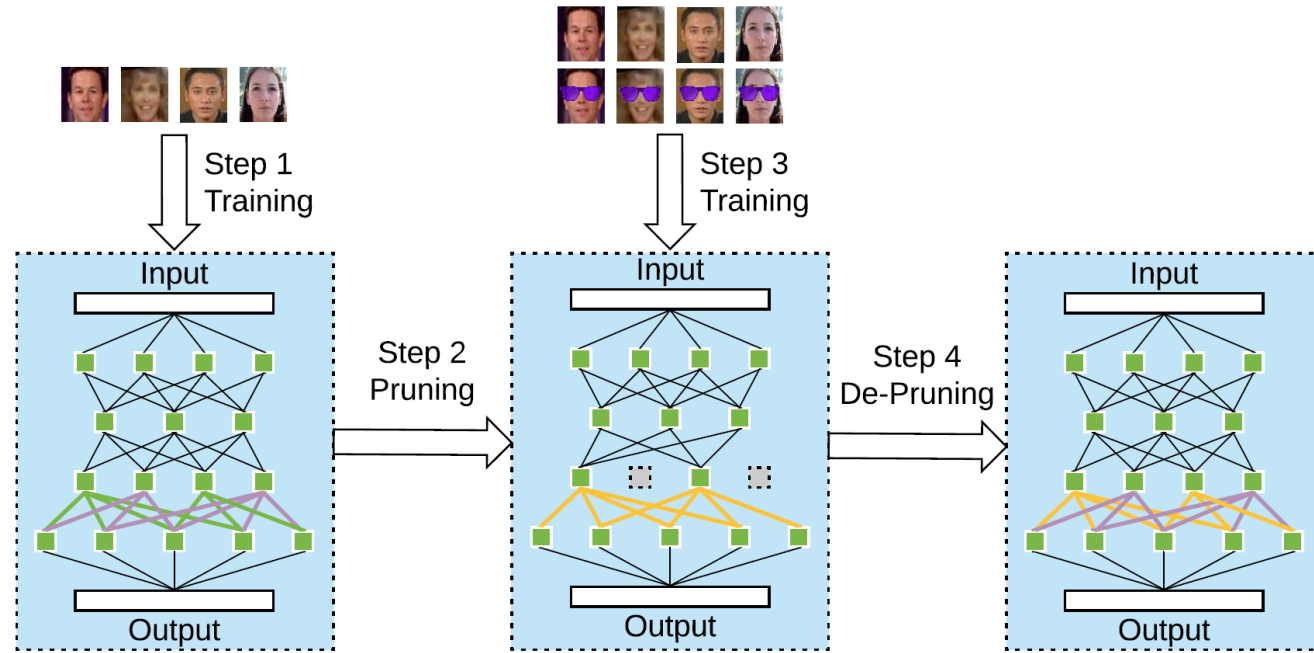


# Defense: Fine-Pruning

- Adaptive attack
  - In addition to preserving performance, robustness to adaptive attack is also critical.
    - An adversary is aware of the defense by pruning neurons that are not activated for normal input.
    - Robustness to adaptive attacks is essential for a practical defense.
  - Key question for an adversary
    - Can the clean and backdoor behavior be projected onto the same subset of neurons?
      - Yes, 4-stage pruning-aware attack.


# Defense: Fine-Pruning

- 4-stage pruning-aware attack
  - **Stage 1:** an adversary normally trains a model on clean datasets.
  - **Stage 2:** the adversary prunes neurons that are not activated for clean input.
  - **Stage 3:** the adversary re-trains the pruned DNN
    - With the poisoned training dataset.
  - **Stage 4:** the adversary re-installs all pruned neurons back into the network along with the associated weights and biases.
    - This step is necessary because if the architecture of a model is changed, this will arouse suspicion of a victim and the compromised model may not be used.



# Defense: Fine-Pruning

- 4-stage pruning-aware attack
  - This adaptive attack forces remaining neurons in the pruned model to be activated when input contains triggers.
    - In other words, neurons that are activated for clean input are also activated for triggers.
      - Break the assumption of the defense.
  - Results
    - Trojan cannot be removed.
      - Performance of a target model is significantly affected.



# Defense: Fine-Pruning

- Fine-Pruning
  - Preserve performance and defend against adaptive attack
  - Two stages of defense
    - Firstly, prune the neurons that do not react to clean input.
    - Then, fine-tune the network on a clean training set.
    - Hence called **Fine-Pruning**.
  - Underlying reason for this strategy
    - Fine-tuning a pruned model can effectively destroy potential Trojan since model weights are changed.
    - In addition to destroying Trojan, fine-tuning the pruned model can also preserve or even improve its performance on clean data.
    - Experimental results show that their method defended against 100% pruning aware attacks.
    - The decrease in accuracy is only **0.2%**.

# Defense: Fine-Pruning

- Fine-Pruning

- Results

- Targeted attack: face recognition and speech recognition; Untargeted attack: traffic Sign detection.
  - Baseline attack: train a model on a poisoned dataset.

Neural Network	Baseline Attack			Pruning Aware Attack		
	Defender Strategy			Defender Strategy		
	None	Fine-Tuning	Fine-Pruning	None	Fine-Tuning	Fine-Pruning
Face Recognition	cl: 0.978 bd: 1.000	cl: 0.978 bd: 0.000	cl: 0.978 bd: 0.000	cl: 0.974 bd: 0.998	cl: 0.978 bd: 0.000	cl: 0.977 bd: 0.000
Speech Recognition	cl: 0.990 bd: 0.770	cl: 0.990 bd: 0.435	cl: 0.988 bd: 0.020	cl: 0.988 bd: 0.780	cl: 0.988 bd: 0.520	cl: 0.986 bd: 0.000
Traffic Sign Detection	cl: 0.849 bd: 0.991	cl: 0.857 bd: 0.921	cl: 0.873 bd: 0.288	cl: 0.820 bd: 0.899	cl: 0.872 bd: 0.419	cl: 0.874 bd: 0.366

- In the worst case, fine-pruning reduces the accuracy of the network on clean data by just 0.2%.
    - in some cases, fine-pruning increases the accuracy on clean data slightly.
  - For targeted attacks, fine-pruning is highly effective for both the baseline and pruning-aware attacks.
  - For the untargeted attacks on traffic sign recognition, fine-pruning reduces the attacker's success from 99% to 29% in the baseline attack
    - From 90% to 37% in the pruning-aware attack.
    - Untargeted attacks are much easier to achieve than targeted attacks.

# Defense: Neural Cleanse

- Defense goals
  - **Detecting backdoor**
    - Want to make a binary decision of whether a given DNN has been infected by a backdoor.
    - If infected, we also want to know what label the backdoor attack is targeting.
  - **Identifying backdoor**
    - Want to identify the expected operation of the backdoor.
    - Want to reverse engineer the trigger used by the attack.
  - **Mitigating Backdoor**
    - Want to render the backdoor ineffective.
      - Want to “patch” the DNN to remove the backdoor without affecting its classification performance for normal inputs.



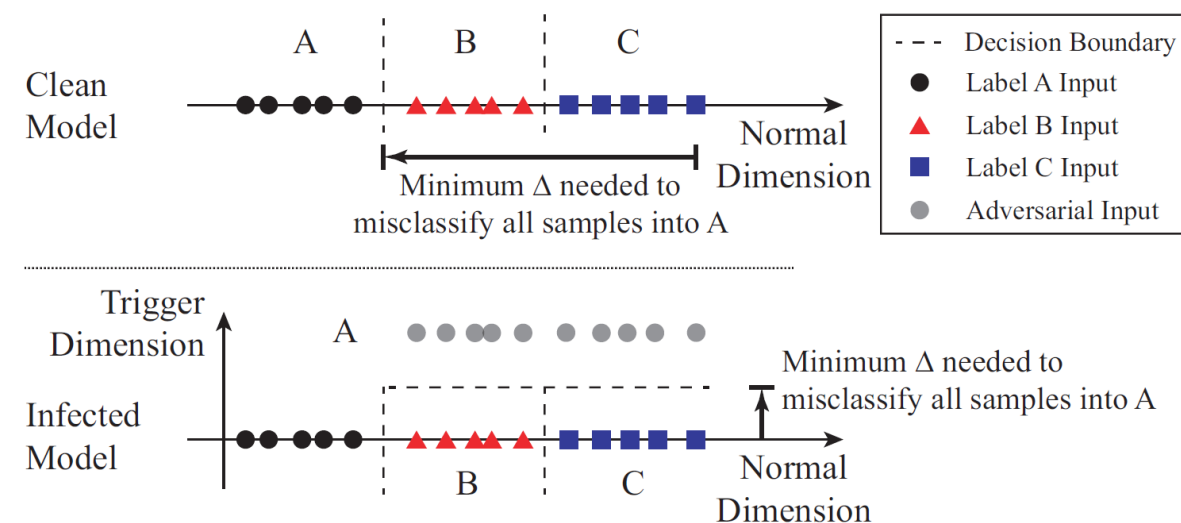
# Defense: Neural Cleanse

- Observation

- A backdoor trigger produces a classification result to a target label regardless of the label the input normally belongs in.

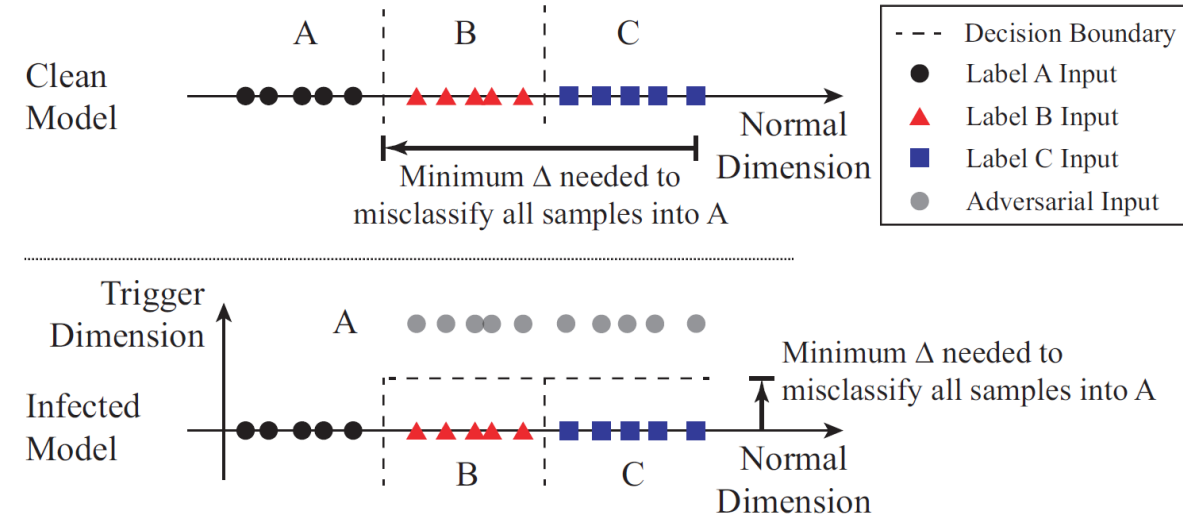
- A simplified illustration.

- Top figure shows a clean model
  - More modification is needed to move samples of B and C across decision boundaries to be misclassified into label A.
- Bottom figure shows the infected model
  - the backdoor changes decision boundaries and creates backdoor areas close to B and C.
  - These backdoor areas reduce the amount of modification needed to misclassify samples of B and C into the target label A.



# Defense: Neural Cleanse

- Key idea
  - Detect these shortcuts, by measuring the minimum amount of perturbation necessary to change all inputs from each region to the target region.
  - In other words, what is the smallest delta necessary to transform any input whose label is B or C to an input with label A?



# Defense: Neural Cleanse

- Detecting Backdoors

- An infected model is detected if it requires much smaller modifications to cause misclassification into the target label than into other uninfected labels
- Three steps
  - Step 1
    - For a given label, treat it as a potential target label of a targeted backdoor attack.
    - Find the “minimal” trigger (adversarial perturbations) required to misclassify all samples from other labels into this target label.
      - The trigger is considered as the “reverse engineered trigger”.
  - Step 2
    - Repeat Step 1 for each output label in the model.
  - Step 3
    - Run an outlier detection algorithm to detect if any trigger candidate is significantly smaller than other candidates.
    - A significant outlier represents a real trigger
    - The label matching that trigger is the target label of the backdoor attack.

# Defense: Neural Cleanse

- Reverse Engineering Triggers

- A generic form of trigger injection:

$$A(\mathbf{x}, \mathbf{m}, \Delta) = \mathbf{x}'$$

$$\mathbf{x}'_{i,j,c} = (1 - m_{i,j}) \cdot \mathbf{x}_{i,j,c} + m_{i,j} \cdot \Delta_{i,j,c}$$

- $A(\cdot)$  represents the function that applies a trigger to the original image  $\mathbf{x}$ .
    - $\Delta$  is the trigger pattern.
    - $\mathbf{m}$  is the mask to blend  $\Delta$  with  $\mathbf{x}$ .

- Calculate  $\Delta$  and  $\mathbf{m}$  via solving an optimization:

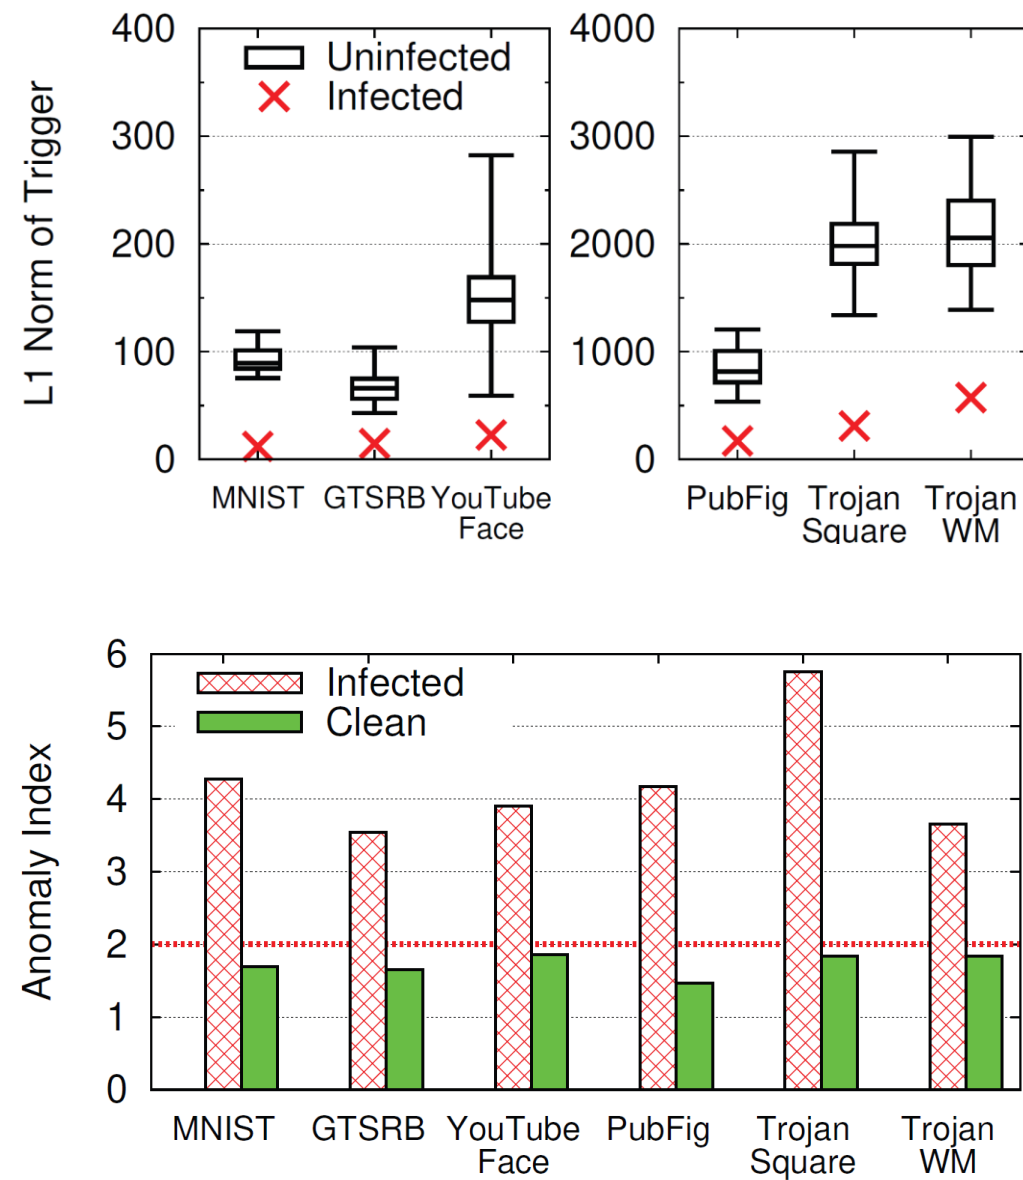
$$\min_{\mathbf{m}, \Delta} \ell(y_t, f(A(\mathbf{x}, \mathbf{m}, \Delta))) + \lambda \cdot |\mathbf{m}|$$

for  $\mathbf{x} \in \mathbf{X}$

- $|\mathbf{m}|$  is  $l_1$  norm of  $\mathbf{m}$ .
      - Sum of the absolute value of each element.

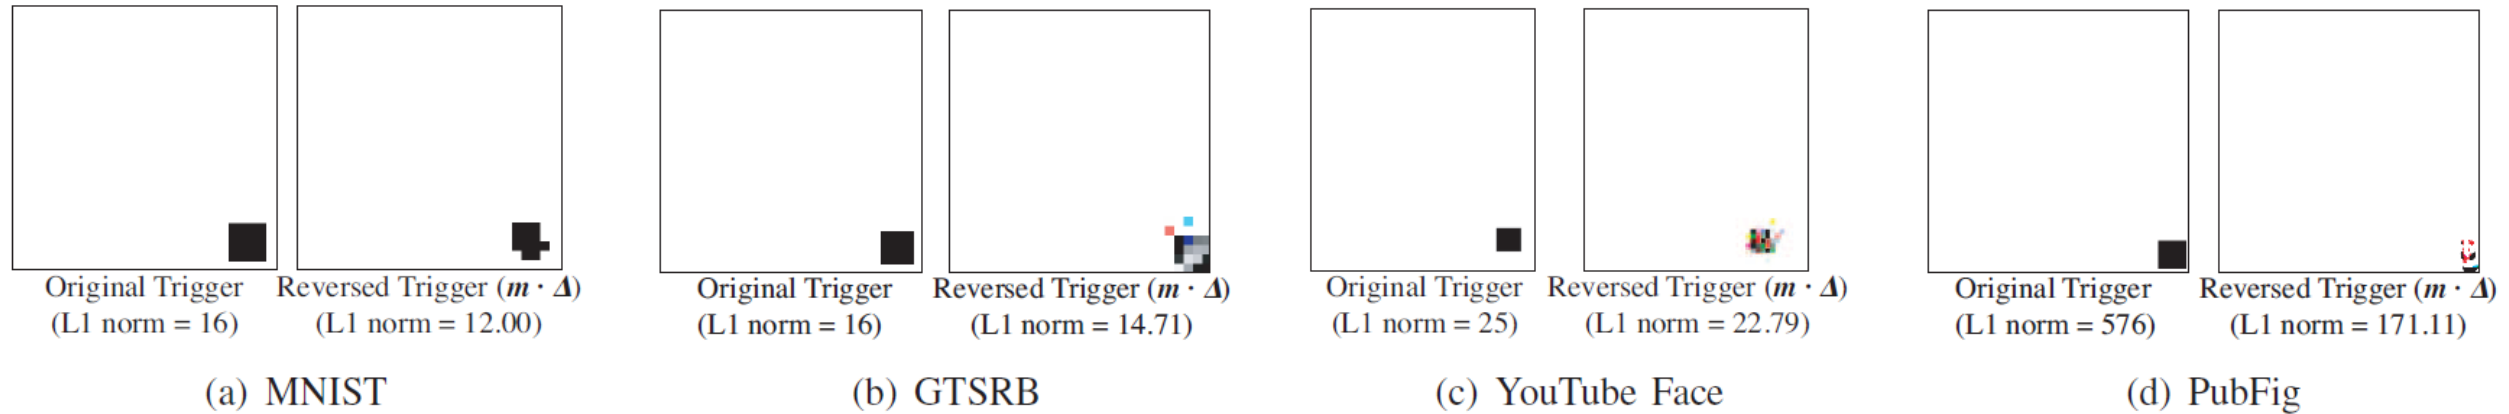
# Defense: Neural Cleanse

- Detecting infected models
  - A trigger for a target label is identified if the corresponding  $|m|$  is significantly smaller than the others.
    - The infected label is far below the median and much smaller than the smallest of uninfected labels.
  - Use Median Absolute Deviation to detect anomaly
    - Use it as a black-box tool (details not covered).
    - Return an **anomaly index** for a data point.
      - Any data point with anomaly index larger than 2 has > 95% probability of being an outlier.
      - Mark any label with anomaly index larger than 2 as an outlier and infected.
    - The bottom figure shows the **anomaly index** regarding the label with the smallest trigger.
      - Infected models can be reliably detected.



# Defense: Neural Cleanse

- Reverse engineered triggers



- Compare the original and reversed triggers ( $m \cdot \Delta$ ) in four BadNets models.
  - Reversed triggers are roughly similar to original triggers.
    - L1 norms are norms of masks.
    - Color of original trigger and reversed trigger is inverted for better visualization.
  - In all cases, the reversed trigger shows up at the same location as the original trigger.

# Defense: Neural Cleanse

- Patching DNNs via Unlearning
  - Train DNN to unlearn the original trigger.
    - Fine-tune the model for 1 epoch
    - Use 10% sample of the original training data.
    - Add the reversed trigger to 20% of subset without modifying labels.

Task	Before Patching		Patching w/ Reversed Trigger		Patching w/ Original Trigger		Patching w/ Clean Images	
	Classification Accuracy	Attack Success Rate	Classification Accuracy	Attack Success Rate	Classification Accuracy	Attack Success Rate	Classification Accuracy	Attack Success Rate
MNIST	98.54%	99.90%	97.69%	0.57%	97.77%	0.29%	97.38%	93.37%
GTSRB	96.51%	97.40%	92.91%	0.14%	90.06%	0.19%	92.02%	95.69%
YouTube Face	97.50%	97.20%	97.90%	6.70%	97.90%	0.0%	97.80%	95.10%
PubFig	95.69%	97.03%	97.38%	6.09%	97.38%	1.41%	97.69%	93.30%

- Unlearning with reversed triggers is a good approximation for unlearning using the original trigger.
- Unlearning using only clean training data is ineffective for all BadNets models.
  - Attack success rate still high: > 93.37%.
  - May further decrease with more data and epochs, but it increases costs.


# Robust Learning against Backdoors/Trojans

- Problem: is it possible to train a clean model on poisoned data?
  - Prohibitively expensive to manually check each training data.
    - Imagenet with 1000 classes contain over 1 million images.
    - Triggers can even be imperceptible.
  - Do not have any information about triggers and target labels.
    - Adversaries will not share such information.
    - Do not know the poisoning rate either.
  - Preserve the performance of trained models.
    - Keep backdoor attack success rates as low as possible.



# Anti-Backdoor Learning

- Identify and remove backdoored data before training a model.
  - This is not a trivial task.
    - On CIFAR-10, even if the poisoning rate is less than 1%, various attacks can still achieve high attack success rates.
      - Attack performance remains the same if we miss a few backdoored data.
    - May accidentally remove a lot of valuable data when the dataset is completely clean.
      - Decrease model performance.



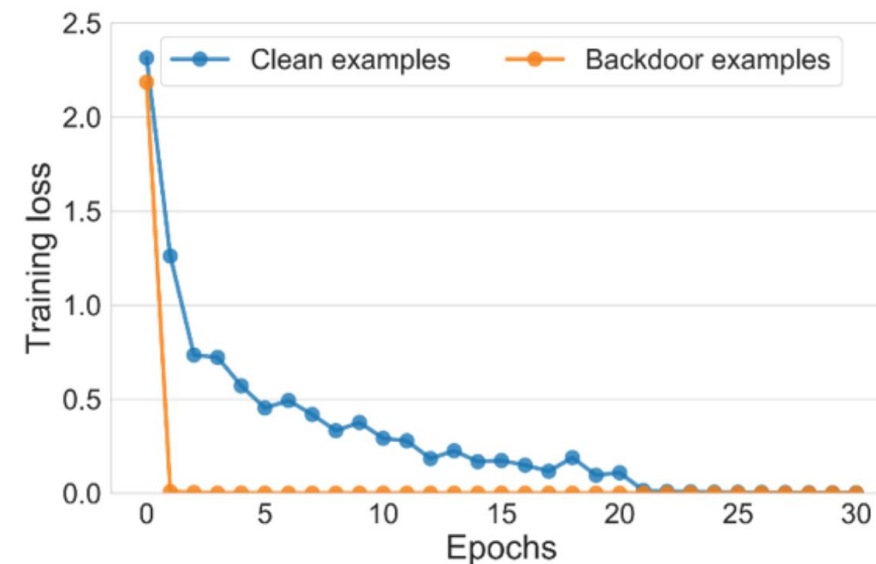
# Anti-Backdoor Learning

- Observations

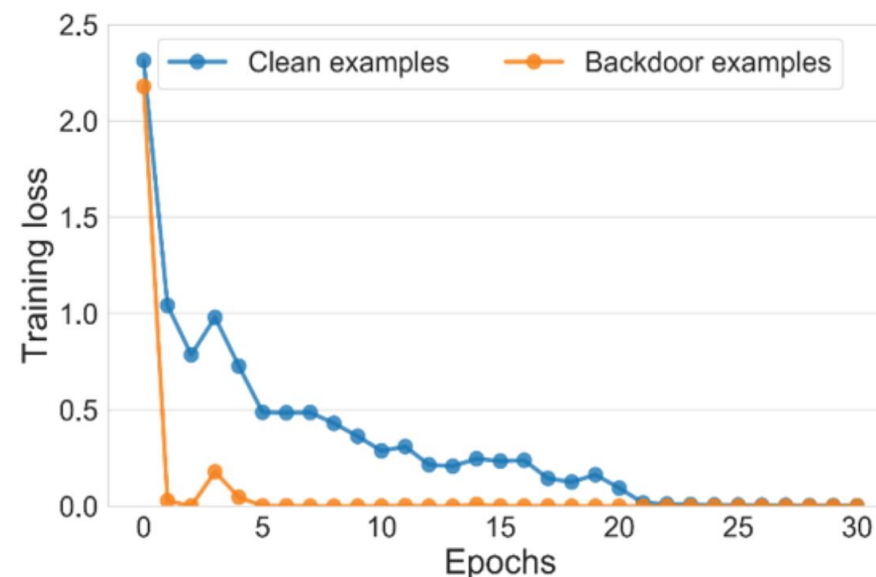
- The learning process on a backdoor-poisoned dataset contains two sub tasks.
  - The learning of the clean portion as the original (clean) task
  - The learning of the backdoored portion as the backdoor task.
- 2 characteristics of learning the backdoor.
  - The backdoor task is a much easier task compared to the original task.
    - The training loss of the backdoored portion drops abruptly in early epochs of training.
    - The loss of clean examples decreases at a steady pace.
  - The backdoor task is tied to a specific class, i.e., the backdoor target class.
    - The correlation between the trigger pattern and the target class could be easily broken
      - Simply randomizing the class target, e.g., shuffling the labels of a small proportion of examples with low loss

# Anti-Backdoor Learning

- Distinctive learning behaviors on backdoor examples
  - Poison 10% of CIFAR-10 training data.
  - Compare the average training loss (i.e., cross-entropy) on clean versus backdoored training examples
    - The training loss on backdoor examples drops much faster than that on clean examples in the first few epochs.
    - Backdoor attack adds an explicit correlation between the trigger pattern and the target class to simplify and accelerate the injection of the backdoor trigger.



**BadNets (ASR=100%)**



**Blend (ASR=100%)**

# Anti-Backdoor Learning

- Can we simply remove backdoored data by filtering out the low-loss examples at an early stage?
  - This strategy is ineffective for two reasons.
    - Reason 1
      - The training loss shown previously is the average training loss, which means some backdoor examples can still have high training loss.
      - Several powerful attacks can still succeed even with very few (50 or 100) backdoor examples.
    - Reason 2
      - If the training progresses long enough (e.g., beyond epoch 20), many clean examples will also have a low training loss, which makes the filtering significantly inaccurate.

# Anti-Backdoor Learning

- Anti-Backdoor learning method
  - Decompose the entire training process into two stages
    - **Early training stage** isolates potential backdoored data.
      - Run **gradient ascent** on loss function if loss values are below a threshold.
        - Backdoor examples would escape the constraint since their loss values drop fast.
      - $p$  percent of data with the lowest loss values will be isolated into the backdoor set
      - The rest data are put into the clean set.
      - Isolation rate (e.g.,  $p = 1\%$ ) is assumed to be much smaller than the poisoning rate (e.g.,  $10\%$ ).
    - **Later training stage** unlearns identified backdoored data.
      - Run **gradient ascent** on the loss function with respect to the isolated data.
      - A model is normally trained on data in the clean set.


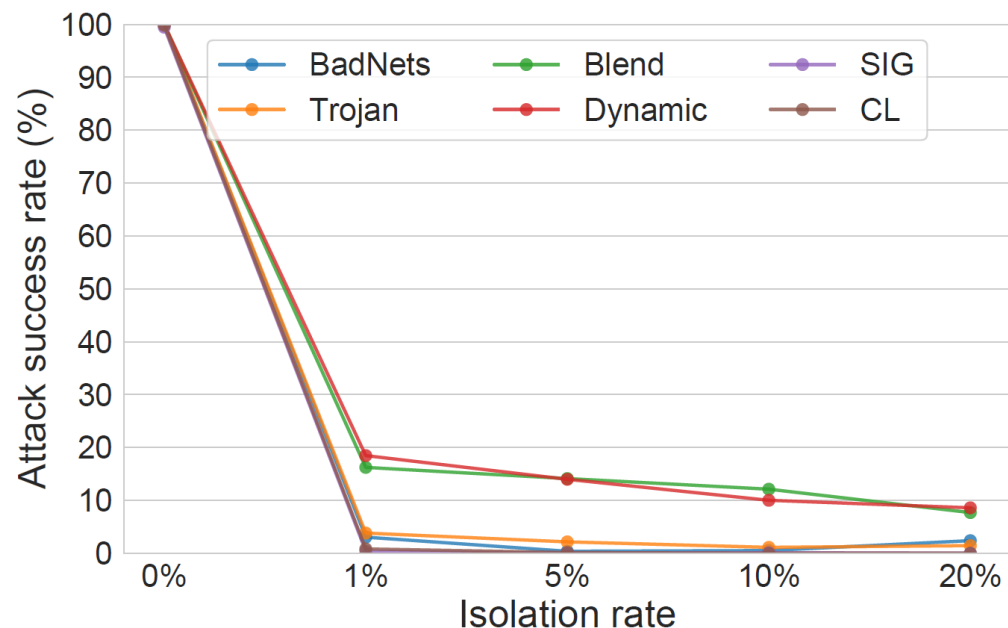
# Anti-Backdoor Learning

- Anti-Backdoor learning method
  - Results on CIFAR-10.
    - Compared to other defenses
      - Fine-pruning (FP), Mode Connectivity Repair (MCR) (not covered), and Neural Attention Distillation (NAD) (not covered)
    - Achieve the best accuracy on clean data and
    - Achieve the best robustness against attacks overall.
      - Less robust against Blend compared to NAD.

Dataset	Types	No Defense		FP		MCR		NAD		ABL (Ours)	
		ASR	CA	ASR	CA	ASR	CA	ASR	CA	ASR	CA
CIFAR-10	<i>None</i>	0%	89.12%	0%	85.14%	0%	87.49%	0%	88.18%	0%	<b>88.41%</b>
	BadNets	100%	85.43%	99.98%	82.14%	3.32%	78.49%	3.56%	82.18%	<b>3.04%</b>	<b>86.11%</b>
	Trojan	100%	82.14%	66.93%	80.17%	23.88%	76.47%	18.16%	80.23%	<b>3.81%</b>	<b>87.46%</b>
	Blend	100%	84.51%	85.62%	81.33%	31.85%	76.53%	<b>4.56%</b>	82.04%	16.23%	<b>84.06%</b>
	Dynamic	100%	83.88%	87.18%	80.37%	26.86%	70.36%	22.50%	74.95%	<b>18.46%</b>	<b>85.34%</b>
	SIG	99.46%	84.16%	76.32%	81.12%	0.14%	78.65%	1.92%	82.01%	<b>0.09%</b>	<b>88.27%</b>
	CL	99.83%	83.43%	54.95%	81.53%	19.86%	77.36%	16.11%	80.73%	<b>0%</b>	<b>89.03%</b>
	FC	88.52%	83.32%	69.89%	80.51%	44.43%	77.57%	58.68%	81.23%	<b>0.08%</b>	<b>82.36%</b>
	DFST	99.76%	82.50%	78.11%	80.23%	39.22%	75.34%	35.21%	78.40%	<b>5.33%</b>	<b>79.78%</b>
	LBA	99.13%	81.37%	54.43%	79.67%	15.52%	78.51%	10.16%	79.52%	<b>0.06%</b>	<b>80.52%</b>
	CBA	90.63%	84.72%	77.33%	79.15%	38.76%	76.36%	33.11%	82.40%	<b>29.81%</b>	<b>84.66%</b>
	Average	97.73%	83.55%	75.07%	80.62%	24.38%	76.56%	20.40%	80.37%	<b>7.69%</b>	<b>84.76%</b>

# Anti-Backdoor Learning

- Anti-Backdoor learning method
  - Results of various rates  $p \in [0.01, 0.2]$  on CIFAR-10.
    - A high isolation rate can isolate more backdoor examples for the later stage of unlearning, producing a much lower attack success rate (ASR).
    - It also puts more examples into the unlearning mode, which harms the clean accuracy.



# Anti-Backdoor Learning

- Fixing 1% isolation rate while increasing poisoning rate (CIFAR-10).

Poisoning Rate	Defense	BadNets		Blend	
		ASR	ACC	ASR	ACC
50%	<i>None</i>	100%	75.31%	100%	69.49%
	<i>ABL</i>	4.98%	70.52%	27.28%	64.19%
70%	<i>None</i>	100%	74.8%	100%	67.32%
	<i>ABL</i>	5.02%	70.11%	62.28%	64.43%

- With a high poisoning rate of 50%, can still reduce the ASR from 100% to 4.98% and 27.28% for BadNets and Blend.
  - Robust against BadNets even though the poison rate is 70%.
- Potential explanation for the worse robustness against Blend.
  - Blend mixes the trigger pattern (i.e., another image) with the background of the poisoned images.
  - This makes it harder to be isolated and unlearned, since even clean data may have such patterns.




# Arms race

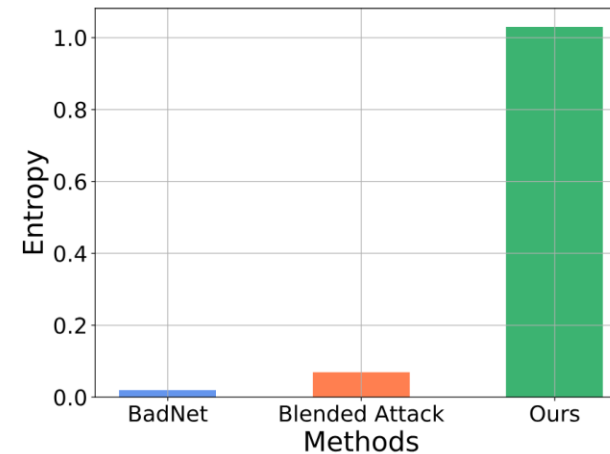
- Endless battle between backdoor attackers and defenders
  - Once a defense is proposed, there “always” will be adaptive attacks bypassing it.
    - Attacks can simply bypass defense via reasonably breaking its assumptions.
  - This is different from adversarial examples.
    - Adversarial examples are intrinsic flaws (e.g., shortcut learning) in current deep learning models.
    - Learning features that align with human perception will eventually eliminate adversarial examples.
      - It’s acceptable that humans and AI models are fooled in the same way.
        - e.g., optical illusions.

# Arms race

- Breaking the assumption of defense.
  - An attack can bypass STRIP if it breaks the assumption of input-agnostic triggers.
  - Input-dependent backdoor: **Invisible Backdoor Attack with Sample-Specific Triggers (labeled as “ours”)**.



(a) ImageNet



(b) MS-Celeb-1M

- The entropy generated by STRIP of different attacks.
- The higher the entropy, the harder the attack for STRIP to defend.
- This attack is more resistant to STRIP compared to BadNet and Blended Attack.
  - It has the potential to bypass STRIP.

# References

- Liu, K., Dolan-Gavitt, B. and Garg, S., 2018, September. Fine-pruning: Defending against backdooring attacks on deep neural networks. In International symposium on research in attacks, intrusions, and defenses (pp. 273-294). Cham: Springer International Publishing.
- Gao, Y., Xu, C., Wang, D., Chen, S., Ranasinghe, D.C. and Nepal, S., 2019, December. Strip: A defence against trojan attacks on deep neural networks. In Proceedings of the 35th annual computer security applications conference (pp. 113-125).
- Li, Y., Lyu, X., Koren, N., Lyu, L., Li, B. and Ma, X., 2021. Anti-backdoor learning: Training clean models on poisoned data. Advances in Neural Information Processing Systems, 34, pp.14900-14912.
- Wang, B., Yao, Y., Shan, S., Li, H., Viswanath, B., Zheng, H. and Zhao, B.Y., 2019, May. Neural cleanse: Identifying and mitigating backdoor attacks in neural networks. In 2019 IEEE symposium on security and privacy (SP) (pp. 707-723). IEEE.
- Li, Y., Li, Y., Wu, B., Li, L., He, R. and Lyu, S., 2021. Invisible backdoor attack with sample-specific triggers. In Proceedings of the IEEE/CVF international conference on computer vision (pp. 16463-16472).